

```
In [9]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn import metrics
```

```
In [3]: import warnings
warnings.filterwarnings('ignore')
print("Warnings ignored!!")
```

Warnings ignored!!

```
In [6]: data=np.load("olivetti_faces.npy")
target=np.load("olivetti_faces_target.npy")
```

```
In [10]: print("There are {} images in the dataset".format(len(data)))
print("There are {} unique targets in the dataset".format(len(np.unique(target))))
print("Size of each image is {}x{}".format(data.shape[1],data.shape[2]))
print("Pixel values were scaled to [0,1] interval. e.g:{}".format(data[0][0,:4]))
```

There are 400 images in the dataset  
There are 40 unique targets in the dataset  
Size of each image is 64x64  
Pixel values were scaled to [0,1] interval. e.g:[0.30991736 0.3677686 0.41735536 0.44214877]

```
In [11]: print("unique target number:",np.unique(target))
```


unique target number: [ 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39]

```
In [12]: def show_40_distinct_people(images, unique_ids):
    #Creating 4X10 subplots in 18x9 figure size
    fig, axarr=plt.subplots(nrows=4, ncols=10, figsize=(18, 9))
    #For easy iteration flattened 4X10 subplots matrix to 40 array
    axarr=axarr.flatten()

    #iterating over user ids
    for unique_id in unique_ids:
        image_index=unique_id*10
        axarr[unique_id].imshow(images[image_index], cmap='gray')
        axarr[unique_id].set_xticks([])
        axarr[unique_id].set_yticks([])
        axarr[unique_id].set_title("face id:{}".format(unique_id))
    plt.suptitle("There are 40 distinct people in the dataset")
```

In [13]: `show_40_distinct_people(data, np.unique(target))`

There are 40 distinct people in the dataset




```
In [14]: def show_10_faces_of_n_subject(images, subject_ids):
    cols=10# each subject has 10 distinct face images
    rows=(len(subject_ids)*10)/cols #
    rows=int(rows)

    fig, axarr=plt.subplots(nrows=rows, ncols=cols, figsize=(18,9))
    #axarr=axarr.flatten()

    for i, subject_id in enumerate(subject_ids):
        for j in range(cols):
            image_index=subject_id*10 + j
            axarr[i,j].imshow(images[image_index], cmap="gray")
            axarr[i,j].set_xticks([])
            axarr[i,j].set_yticks([])
            axarr[i,j].set_title("face id:{}".format(subject_id))
```

In [15]: `show_10_faces_of_n_subject(images=data, subject_ids=[0,5, 21, 24, 36])`



```
In [16]: X=data.reshape((data.shape[0],data.shape[1]*data.shape[2]))
print("X shape:",X.shape)
```


X shape: (400, 4096)

```
In [17]: X_train, X_test, y_train, y_test=train_test_split(X, target, test_size=0.3, stratify=target, random_state=
print("X_train shape:",X_train.shape)
print("y_train shape:{}".format(y_train.shape))
```

X\_train shape: (280, 4096)  
y\_train shape:(280,)

```
In [18]: y_frame=pd.DataFrame()
y_frame['subject ids']=y_train
y_frame.groupby(['subject ids']).size().plot.bar(figsize=(15,8),title="Number of Samples for Each Classes"
```

Out[18]: <AxesSubplot:title={'center':'Number of Samples for Each Classes'}, xlabel='subject ids'>




```
In [20]: from sklearn.decomposition import PCA
pca=PCA(n_components=2)
pca.fit(X)
X_pca=pca.transform(X)
```

```
In [21]: number_of_people=10
index_range=number_of_people*10
fig=plt.figure(figsize=(10,8))
ax=fig.add_subplot(1,1,1)
scatter=ax.scatter(X_pca[:index_range,0],
                  X_pca[:index_range,1],
                  c=target[:index_range],
                  s=10,
                  cmap=plt.get_cmap('jet', number_of_people)
                )


ax.set_xlabel("First Principle Component")
ax.set_ylabel("Second Principle Component")
ax.set_title("PCA projection of {} people".format(number_of_people))

fig.colorbar(scatter)
```

Out[21]: <matplotlib.colorbar.Colorbar at 0x1e2db5ec040>



```
In [22]: pca=PCA()  
pca.fit(X)  
  
plt.figure(1, figsize=(12,8))  
  
plt.plot(pca.explained_variance_, linewidth=2)  
  
plt.xlabel('Components')  
plt.ylabel('Explained Variaces')  
plt.show()
```



```
In [23]: n_components=90
```

```
In [24]: pca=PCA(n_components=n_components, whiten=True)  
pca.fit(X_train)
```

```
Out[24]: PCA(n_components=90, whiten=True)
```

```
In [25]: PCA(copy=True, iterated_power='auto', n_components=90, random_state=None,  
          svd_solver='auto', tol=0.0, whiten=True)
```

```
Out[25]: PCA(n_components=90, whiten=True)
```

```
In [26]: fig,ax=plt.subplots(1,1,figsize=(8,8))
ax.imshow(pca.mean_.reshape((64,64)), cmap="gray")
ax.set_xticks([])
ax.set_yticks([])
ax.set_title('Average Face')
```

Out[26]: Text(0.5, 1.0, 'Average Face')



```

In [27]: number_of_eigenfaces=len(pca.components_)
eigen_faces=pca.components_.reshape((number_of_eigenfaces, data.shape[1], data.shape[2]))

cols=10
rows=int(number_of_eigenfaces/cols)
fig, axarr=plt.subplots(nrows=rows, ncols=cols, figsize=(15,15))
axarr=axarr.flatten()
for i in range(number_of_eigenfaces):
    axarr[i].imshow(eigen_faces[i],cmap="gray")
    axarr[i].set_xticks([])
    axarr[i].set_yticks([])
    axarr[i].set_title("eigen id:{}".format(i))
plt.suptitle("All Eigen Faces".format(10* "=", 10* "="))

```

Out[27]: Text(0.5, 0.98, 'All Eigen Faces')

All Eigen Faces



```

In [28]: X_train_pca=pca.transform(X_train)
X_test_pca=pca.transform(X_test)


```

```
In [29]: clf = SVC()
clf.fit(X_train_pca, y_train)
y_pred = clf.predict(X_test_pca)
print("accuracy score:{:.2f}".format(metrics.accuracy_score(y_test, y_pred)))
```

accuracy score:0.93

```
In [30]: import seaborn as sns
plt.figure(1, figsize=(12,8))
sns.heatmap(metrics.confusion_matrix(y_test, y_pred))
```

Out[30]: <AxesSubplot:>





```
In [31]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.50	0.33	0.40	3
1	1.00	1.00	1.00	3
2	0.67	0.67	0.67	3
3	1.00	1.00	1.00	3
4	1.00	0.67	0.80	3
5	1.00	1.00	1.00	3
6	1.00	1.00	1.00	3
7	1.00	0.67	0.80	3
8	1.00	1.00	1.00	3
9	1.00	1.00	1.00	3
10	1.00	1.00	1.00	3
11	1.00	1.00	1.00	3
12	1.00	0.67	0.80	3
13	1.00	1.00	1.00	3
14	1.00	1.00	1.00	3
15	0.75	1.00	0.86	3
16	1.00	1.00	1.00	3
17	1.00	1.00	1.00	3
18	1.00	1.00	1.00	3
19	1.00	1.00	1.00	3
20	0.60	1.00	0.75	3
21	1.00	0.67	0.80	3
22	1.00	1.00	1.00	3
23	1.00	1.00	1.00	3
24	1.00	1.00	1.00	3
25	1.00	0.67	0.80	3
26	1.00	1.00	1.00	3
27	1.00	1.00	1.00	3
28	1.00	1.00	1.00	3
29	1.00	1.00	1.00	3
30	1.00	1.00	1.00	3
31	1.00	0.67	0.80	3
32	1.00	1.00	1.00	3
33	1.00	1.00	1.00	3
34	1.00	1.00	1.00	3
35	1.00	1.00	1.00	3
36	1.00	1.00	1.00	3
37	1.00	1.00	1.00	3
38	0.75	1.00	0.86	3
39	0.50	1.00	0.67	3
accuracy			0.93	120
macro avg	0.94	0.93	0.92	120
weighted avg	0.94	0.93	0.92	120

```
In [32]: models=[]
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('LR', LogisticRegression()))
models.append(('NB', GaussianNB()))
models.append(('KNN', KNeighborsClassifier(n_neighbors=5)))
models.append(('DT', DecisionTreeClassifier()))
models.append(('SVM', SVC()))

for name, model in models:

    clf=model

    clf.fit(X_train_pca, y_train)

    y_pred=clf.predict(X_test_pca)
    print(10*"=", "{} Result".format(name).upper(), 10*"=")
    print("Accuracy score:{:.2f}".format(metrics.accuracy_score(y_test, y_pred)))
    print()

===== LDA RESULT =====
Accuracy score:0.93

===== LR RESULT =====
Accuracy score:0.93

===== NB RESULT =====
Accuracy score:0.87

===== KNN RESULT =====
Accuracy score:0.70

===== DT RESULT =====
Accuracy score:0.61

===== SVM RESULT =====
Accuracy score:0.93
```

```
In [33]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
pca=PCA(n_components=n_components, whiten=True)
pca.fit(X)
X_pca=pca.transform(X)
for name, model in models:
    kfold=KFold(n_splits=5, shuffle=True, random_state=0)

    cv_scores=cross_val_score(model, X_pca, target, cv=kfold)
    print("{} mean cross validations score:{:.2f}".format(name, cv_scores.mean()))

LDA mean cross validations score:0.98
LR mean cross validations score:0.93
NB mean cross validations score:0.80
KNN mean cross validations score:0.68
DT mean cross validations score:0.46
SVM mean cross validations score:0.87
```


```
In [34]: lr=LinearDiscriminantAnalysis()
lr.fit(X_train_pca, y_train)
y_pred=lr.predict(X_test_pca)
print("Accuracy score:{:.2f}".format(metrics.accuracy_score(y_test, y_pred)))

Accuracy score:0.93
```

```
In [35]: cm=metrics.confusion_matrix(y_test, y_pred)
```

```
plt.subplots(1, figsize=(12,12))  
sns.heatmap(cm)
```

Out[35]: <AxesSubplot:>



```
In [36]: print("Classification Results:\n{}".format(metrics.classification_report(y_test, y_pred)))
```

```
Classification Results:
              precision    recall  f1-score   support

     0           1.00      0.33      0.50         3
     1           1.00      1.00      1.00         3
     2           1.00      1.00      1.00         3
     3           1.00      1.00      1.00         3
     4           0.67      0.67      0.67         3
     5           1.00      1.00      1.00         3
     6           1.00      1.00      1.00         3
     7           1.00      0.67      0.80         3
     8           1.00      1.00      1.00         3
     9           1.00      1.00      1.00         3
    10           1.00      1.00      1.00         3
    11           1.00      1.00      1.00         3
    12           0.50      0.33      0.40         3
    13           1.00      1.00      1.00         3
    14           1.00      1.00      1.00         3
    15           1.00      1.00      1.00         3
    16           1.00      1.00      1.00         3
    17           1.00      1.00      1.00         3
    18           1.00      1.00      1.00         3
    19           1.00      1.00      1.00         3
    20           0.75      1.00      0.86         3
    21           1.00      0.67      0.80         3
    22           0.75      1.00      0.86         3
    23           1.00      1.00      1.00         3
    24           0.60      1.00      0.75         3
    25           1.00      1.00      1.00         3
    26           1.00      0.67      0.80         3
    27           1.00      1.00      1.00         3
    28           1.00      1.00      1.00         3
    29           1.00      1.00      1.00         3
    30           1.00      1.00      1.00         3
    31           1.00      0.67      0.80         3
    32           1.00      1.00      1.00         3
    33           1.00      1.00      1.00         3
    34           1.00      1.00      1.00         3
    35           1.00      1.00      1.00         3
    36           1.00      1.00      1.00         3
    37           1.00      1.00      1.00         3
    38           0.75      1.00      0.86         3
    39           0.60      1.00      0.75         3

 accuracy          0.93          120
 macro avg         0.94          0.93          0.92          120
 weighted avg      0.94          0.93          0.92          120
```

```
In [37]: from sklearn.model_selection import LeaveOneOut
loo_cv=LeaveOneOut()
clf=LogisticRegression()
cv_scores=cross_val_score(clf,
                          X_pca,
                          target,
                          cv=loo_cv)
print("{} Leave One Out cross-validation mean accuracy score:{:.2f}".format(clf.__class__.__name__,
                                                                              cv_scores.mean()))
```

LogisticRegression Leave One Out cross-validation mean accuracy score:0.94

```
In [38]: from sklearn.model_selection import LeaveOneOut
loo_cv=LeaveOneOut()
clf=LinearDiscriminantAnalysis()
cv_scores=cross_val_score(clf,
                          X_pca,
                          target,
                          cv=loo_cv)
print("{} Leave One Out cross-validation mean accuracy score:{:.2f}".format(clf.__class__.__name__,
                                                                              cv_scores.mean()))
```

LinearDiscriminantAnalysis Leave One Out cross-validation mean accuracy score:0.98

```
In [39]: from sklearn.model_selection import GridSearchCV
```

```
In [42]: lr=LogisticRegression(C=1.0, penalty="l2")
lr.fit(X_train_pca, y_train)
print("lr score:{:.2f}".format(lr.score(X_test_pca, y_test)))

lr score:0.93
```

```
In [43]: from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier

Target=label_binarize(target, classes=range(40))
print(Target.shape)
print(Target[0])

n_classes=Target.shape[1]

(400, 40)
[1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0]
```

```
In [44]: X_train_multiclass, X_test_multiclass, y_train_multiclass, y_test_multiclass=train_test_split(X,
                                                                                                     Target,
                                                                                                     test_size=0.3
                                                                                                     stratify=Target,
                                                                                                     random_state=0)
```

```
In [45]: pca=PCA(n_components=n_components, whiten=True)
pca.fit(X_train_multiclass)

X_train_multiclass_pca=pca.transform(X_train_multiclass)
X_test_multiclass_pca=pca.transform(X_test_multiclass)
```

```
In [46]: oneRestClassifier=OneVsRestClassifier(lr)

oneRestClassifier.fit(X_train_multiclass_pca, y_train_multiclass)
y_score=oneRestClassifier.decision_function(X_test_multiclass_pca)
```

```
In [47]: # For each class
precision = dict()
recall = dict()
average_precision = dict()
for i in range(n_classes):
    precision[i], recall[i], _ = metrics.precision_recall_curve(y_test_multiclass[:, i],
                                                                y_score[:, i])
    average_precision[i] = metrics.average_precision_score(y_test_multiclass[:, i], y_score[:, i])

# A "micro-average": quantifying score on all classes jointly
precision["micro"], recall["micro"], _ = metrics.precision_recall_curve(y_test_multiclass.ravel(),
                                                                y_score.ravel())
average_precision["micro"] = metrics.average_precision_score(y_test_multiclass, y_score,
                                                                average="micro")
print('Average precision score, micro-averaged over all classes: {0:0.2f}'
      .format(average_precision["micro"]))
```

Average precision score, micro-averaged over all classes: 0.97

```
In [49]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
In [52]: print("Accuracy score:{:.2f}".format(metrics.accuracy_score(y_test, y_pred)))
print("Classification Results:\n{}".format(metrics.classification_report(y_test, y_pred)))
```

Accuracy score:0.93

Classification Results:

	precision	recall	f1-score	support
0	1.00	0.33	0.50	3
1	1.00	1.00	1.00	3
2	1.00	1.00	1.00	3
3	1.00	1.00	1.00	3
4	0.67	0.67	0.67	3
5	1.00	1.00	1.00	3
6	1.00	1.00	1.00	3
7	1.00	0.67	0.80	3
8	1.00	1.00	1.00	3
9	1.00	1.00	1.00	3
10	1.00	1.00	1.00	3
11	1.00	1.00	1.00	3
12	0.50	0.33	0.40	3
13	1.00	1.00	1.00	3
14	1.00	1.00	1.00	3
15	1.00	1.00	1.00	3
16	1.00	1.00	1.00	3
17	1.00	1.00	1.00	3
18	1.00	1.00	1.00	3
19	1.00	1.00	1.00	3
20	0.75	1.00	0.86	3
21	1.00	0.67	0.80	3
22	0.75	1.00	0.86	3
23	1.00	1.00	1.00	3
24	0.60	1.00	0.75	3
25	1.00	1.00	1.00	3
26	1.00	0.67	0.80	3
27	1.00	1.00	1.00	3
28	1.00	1.00	1.00	3
29	1.00	1.00	1.00	3
30	1.00	1.00	1.00	3
31	1.00	0.67	0.80	3
32	1.00	1.00	1.00	3
33	1.00	1.00	1.00	3
34	1.00	1.00	1.00	3
35	1.00	1.00	1.00	3
36	1.00	1.00	1.00	3
37	1.00	1.00	1.00	3
38	0.75	1.00	0.86	3
39	0.60	1.00	0.75	3
accuracy			0.93	120
macro avg	0.94	0.93	0.92	120
weighted avg	0.94	0.93	0.92	120

```
In [53]: from sklearn.pipeline import Pipeline
```

```
In [55]: print("Accuracy score:{:.2f}".format(metrics.accuracy_score(y_test, y_pred)))
print("Classification Results:\n{}".format(metrics.classification_report(y_test, y_pred)))
```

Accuracy score:0.93

Classification Results:

	precision	recall	f1-score	support
0	1.00	0.33	0.50	3
1	1.00	1.00	1.00	3
2	1.00	1.00	1.00	3
3	1.00	1.00	1.00	3
4	0.67	0.67	0.67	3
5	1.00	1.00	1.00	3
6	1.00	1.00	1.00	3
7	1.00	0.67	0.80	3
8	1.00	1.00	1.00	3
9	1.00	1.00	1.00	3
10	1.00	1.00	1.00	3
11	1.00	1.00	1.00	3
12	0.50	0.33	0.40	3
13	1.00	1.00	1.00	3
14	1.00	1.00	1.00	3
15	1.00	1.00	1.00	3
16	1.00	1.00	1.00	3
17	1.00	1.00	1.00	3
18	1.00	1.00	1.00	3
19	1.00	1.00	1.00	3
20	0.75	1.00	0.86	3
21	1.00	0.67	0.80	3
22	0.75	1.00	0.86	3
23	1.00	1.00	1.00	3
24	0.60	1.00	0.75	3
25	1.00	1.00	1.00	3
26	1.00	0.67	0.80	3
27	1.00	1.00	1.00	3
28	1.00	1.00	1.00	3
29	1.00	1.00	1.00	3
30	1.00	1.00	1.00	3
31	1.00	0.67	0.80	3
32	1.00	1.00	1.00	3
33	1.00	1.00	1.00	3
34	1.00	1.00	1.00	3
35	1.00	1.00	1.00	3
36	1.00	1.00	1.00	3
37	1.00	1.00	1.00	3
38	0.75	1.00	0.86	3
39	0.60	1.00	0.75	3
accuracy			0.93	120
macro avg	0.94	0.93	0.92	120
weighted avg	0.94	0.93	0.92	120